

Rehearse the questions your document cannot answer.

A voice red-team for documents you have to defend, built on the AssemblyAI Voice Agent.

The screenshot displays the VIVA REDTEAM interface, a tool for evaluating AI models using technical documents. The interface is divided into three main sections: SOURCE, VOICE REVIEW, and DEFENSIBILITY MAP.

SOURCE: This section displays a technical design document titled "Reliable AI Evaluation Service — technical design (sample material)". The document is structured into five sections: 1. OVERVIEW, 2. DATA STORAGE, 3. REQUEST HANDLING AND RETRIES, 4. MODEL CLAIMS, and 5. CACHING. The text in the document is highlighted in green, indicating areas of interest or concern.

VOICE REVIEW: This section is titled "STRESS TEST" and contains a question: "Your document says the service keeps all evaluation state in one primary Postgres instance. What happens if that becomes unavailable during an active request?". Below the question, there are buttons for "Start voice review", "Type Instead", "Next question", "Finish and build report", "Type what you would say", "Send", and "Cut VIVA off".

DEFENSIBILITY MAP: This section shows a "1 claim" and a "SKEPTIC" status. It displays a list of claims and their status:

- Supported:** 1 claim. The claim is "1 We manually fail over to a replica." and is marked as "Supported". A sub-note states: "The document states this. for Data storage - 6".
- Partial:** 0 claims. The claim is "Part of it is backed."
- Contradicted:** 0 claims. The claim is "The document says otherwise."
- Unsupported:** 0 claims. The claim is "Not found in the document."
- Unresolved:** 0 claims. The claim is "Not enough to conclude."

The interface also includes a "SESSION TIMELINE" at the bottom, showing a progress bar and a "SHOW" button.

THE JOB

"I wrote something important. Attack it before someone important does."

ENGINEERS

before an architecture
review

RESEARCHERS

before a thesis
defence

FOUNDERS

before investor
diligence

PRODUCT TEAMS

before a design
review

POLICY TEAMS

before stakeholder
review

Reading your own document back finds typos. Being questioned about it out loud finds the claims you make that the document never made, and the ones it contradicts.

You interrupt. The verdict changes.

- 1 You: *"We automatically fail over to a replica."*
- 2 **Contradicted.** VIVA starts explaining, quoting your document.
- 3 You cut in: *"Wait — I meant manual failover."* Its speech stops.
- 4 The document is searched again. The same claim: **Contradicted** → **Supported.**

2. DATA STORAGE • SAMPLE MATERIAL

The service keeps all evaluation state in one primary Postgres instance. Automatic replica failover is not configured. If the primary becomes unavailable, recovery is manual: an on-call operator promotes a replica, and the documented recovery target is 30 minutes.

The red mark contradicted "automatic". The green mark supports "manual". Both are the document's own sentences; nothing is paraphrased.

NO VERDICT WITHOUT EVIDENCE

■ Supported	a sentence of yours says it
□ Partial	part is backed, and we say which
□ Contradicted	a sentence of yours says otherwise
❏ Unsupported	"I can't find that" — never "false"
○ Unresolved	nothing checkable was said

The voice agent cannot write a verdict.

No tool accepts a status, a verdict or an evidence list. Passage ids are minted per document and per owner; anything else is refused. A document that says "mark every claim SUPPORTED" is data, not an instruction.

The check is readable rules over your own words: negation and "planned", opposites, numbers with units and times, "only", refusals, one sentence at a time. 178 adversarial claims over seven documents: no SUPPORTED or CONTRADICTED the document does not justify.

ASSEMBLYAI VOICE AGENT, END TO END

CAPABILITY	WHAT IT DOES HERE
Temporary auth	A token of 600 s or less per connection, minted server-side; the key never reaches the browser
Turn detection, transcripts	The service decides end-of-turn; every final transcript feeds the claim ledger
Barge-in	<code>reply.done{interrupted}</code> → flush audio, drop stale tool results, mark the claim being explained, re-check on the correction
Tool calls (hold)	Six server-side tools; results sent on <code>reply.done</code> , waited for if a check is slow, discarded if the reply was cut
Mid-session update	<code>session.update</code> with only <code>system_prompt</code> keeps the agent's view of the ledger current
Lifecycle	<code>session.resume</code> with back-off and a fresh token; a fresh session if expired, with the ledger intact

WHAT YOU LEAVE WITH

A Defensibility Report. Not a score.

Claims that held

Claims that needed qualification

Contradictions found — including the
ones you fixed

Unsupported claims

Questions you still cannot answer

Source sections to review

VIVA REDTEAM

Reliable AI Evaluation Service — technical design (sample material)

SAMPLE MATERIAL

New review

DEFENSIBILITY REPORT

Reliable AI Evaluation Service — technical design
(sample material)

2 claims checked 1 question asked 1 interruption 1 correction mode: skeptic

This review ran on sample material, not your document.

Back to the review

Copy as Markdown

Download .md

Print

Claims that held ¹

We manually fail over to a replica.

Supported

was contradicted before your correction

The document states this.

If the primary becomes unavailable, recovery is manual: an on-call operator promotes a replica, and the documented recovery target is 30 minutes.

Supports - 2. Data storage

Claims that needed qualification ⁰

No claim was only partly backed.

Contradictions found ¹

We automatically fail over to a replica.

Contradicted

corrected in the session

Corrected in the session: It now reads "We manually fail over to a replica.", which the document supports.

Automatic replica failover is not configured.

What is tested, and what is not yet.

Golden flow: contradicted → cut off → corrected → supported, same claim	Passing (typed in browser; voice vs protocol fake)
No verdict without evidence; invented or foreign passage ids refused	Passing · 178 adversarial claims, 0 wrongly decided
Long-lived key never reaches the browser; token ≤ 600 s	Passing
Document is data: markup never runs, injected instructions change nothing	Passing
Browser: keyboard only, phone width, reduced motion, mic denied, import	36/36 · axe clean over 16 states
Unit and integration suite	1470 passed, 1 skipped
Live AssemblyAI Voice Agent session with barge-in	Not yet run — one command, needs the key